

The table shows the results for the serial and parallel version of KNN.

	# of threads	Serial	1	2	4	8	2048
Virtual Machine	Small	36	35 (1.02)	26 (1.38)	20 (1.8)	30 (1.2)	556 (0.06)
	Medium	12810	13342 (0.96)	6273 (2.04)	4212 (3.04)	4798 (2.67)	
Maple Server	Small	61	64 (0.95)	37 (1.64)	20 (3.05)	11 (5.5)	128 (0.47)
	Medium	9395	9589 (0.98)	4863 (1.93)	2529 (3.71)	1378 (6.96)	1995 (4.70)

On the CUDA platform, the result was:

- 226 ms for the small dataset. (speedup = 0.26)
- 215 ms for the medium dataset (speedup = 44).

The results shows that the delay is mostly the overhead of communication with GPU and actual computation for the two dataset doesn't have significant difference. Which probably means that the higher speedups can be achieved with databases of larger size.